

ENVIRONMENT OVERVIEW

Deployment Architecture

This guide deploys a FIPS-compliant OpenShift 4.18 cluster with full Platform Plus capabilities in a completely disconnected air-gapped environment. The deployment includes OpenShift Virtualization for running VMs alongside containers.

Cluster Specifications:

- **Cluster Name:** ovetest
- **Base Domain:** test.com
- **OpenShift Version:** 4.18
- **Network:** 192.168.1.0/24
- **Deployment Type:** Fully disconnected air-gapped with FIPS compliance
- **Features:** OpenShift Virtualization, Advanced Cluster Management, OpenShift Data Foundation

Infrastructure Components:

- 3x Bare metal OpenShift master nodes (host01-03)
- 1x NFS storage server (dedicated bare metal)
- 6x Supporting VMs on vSphere cluster
- Single 1GbE network infrastructure

IP Address Allocation

Supporting Infrastructure

192.168.1.10	tmpregistry.test.com	(Mirror Registry)
192.168.1.11	octools.test.com	(OpenShift Tools - FIPS enabled)
192.168.1.13	dns.test.com	(DNS/NTP Server)
192.168.1.13	content.test.com	(Content Server)
192.168.1.15	nfs.test.com	(NFS Storage Server)

OpenShift Infrastructure

192.168.1.21	host01.ove.test.com	(Master Node 1)
192.168.1.22	host02.ove.test.com	(Master Node 2)
192.168.1.23	host03.ove.test.com	(Master Node 3)
192.168.1.25	bootstrap.ove.test.com	(Bootstrap Node)

OpenShift Services

192.168.1.30	api.ove.test.com	(API VIP)
192.168.1.30	api-int.ove.test.com	(Internal API VIP)
192.168.1.31	*.apps.ove.test.com	(Applications VIP)

PREREQUISITES AND PLANNING

Required Expertise

- OpenShift 4.18 administration and troubleshooting
- RHEL system administration and FIPS compliance
- DNS/NTP service configuration
- Container registry management and mirroring
- Network infrastructure management

Critical FIPS Requirements

⚠ **MANDATORY:** FIPS compliance must be configured from the beginning - it cannot be enabled after installation.

- **Installation Host:** The octools VM must be FIPS-enabled before running any OpenShift installation commands
- **All Systems:** RHEL systems should be FIPS-compliant for consistency
- **Validation:** FIPS mode verification is required at each phase

Hardware Requirements

Bare Metal Servers (4 total):

- 3x OpenShift Master Nodes: 500GB OS drive, 1TB data drive, single 1GbE NIC
- 1x NFS Storage Server: 3TB SSD storage, single 1GbE NIC

vSphere Cluster Requirements:

- Sufficient resources for 6 VMs (total: 28 vCPU, 68GB RAM, 2.3TB storage)
- Network connectivity to 192.168.1.0/24

Pre-Deployment Validation Checklist

- ✓ All bare metal servers have been wiped clean
- ✓ Network infrastructure (192.168.1.0/24) is operational
- ✓ vSphere cluster has sufficient resources allocated
- ✓ External connectivity available for initial image pulling
- ✓ DNS domain (test.com) is available for use
- ✓ Red Hat pull secret downloaded and accessible

PHASE 1: EXTERNAL INFRASTRUCTURE SETUP (INTERNET CONNECTED)

Context: This phase is performed on systems WITH internet access to download and prepare all required OpenShift images and content for transfer to the air-gapped environment.

Step 1.1: Create External Mirror Host (OUTSIDE AIR GAP)

Purpose: Download OpenShift images and operator content from Red Hat registries

Create VM with: 4c, 16GB, 1TB storage, RHEL9, WITH INTERNET ACCESS

Install required packages for image mirroring

```
sudo dnf install -y podman curl wget jq
```

Download oc-mirror v2 tool

```
curl -L https://console.redhat.com/openshift/downloads/tool-mirror-registry -o oc-mirror.tar.gz
```

```
tar xvf oc-mirror.tar.gz
```

```
sudo mv oc-mirror /usr/local/bin/
```

```
chmod +x /usr/local/bin/oc-mirror
```

Verify oc-mirror installation

```
oc-mirror --v2 --help
```

Create working directory for mirror operations

```
mkdir -p ~/mirror-workspace
```

```
cd ~/mirror-workspace
```

Step 1.2: Configure Image Mirroring Content

Purpose: Define what OpenShift content to mirror for the disconnected installation

Download your Red Hat pull secret from <https://console.redhat.com/openshift/install/pull-secret>

Save it as pull-secret.json in current directory

Create comprehensive ImageSetConfiguration for Platform Plus

```
cat << 'EOF' > imageset-config.yaml
```

```
kind: ImageSetConfiguration
```

```
apiVersion: mirror.openshift.io/v2alpha1
```

```
mirror:
```

```
  platform:
```

```
    channels:
```

```
      - name: stable-4.18
```

```
        minVersion: 4.18.0
```

```
        maxVersion: 4.18.20
```

```
        type: ocp
```

```
    graph: true
```

```
    kubeVirtContainer: true
```

```
  additionalImages:
```

```
    # Base Container Images
```

```
    - name: registry.redhat.io/ubi8/ubi:latest
```

```
    - name: registry.redhat.io/ubi9/ubi:latest
```

```
    # RHEL Guest Images for Virtualization
```

```
    - name: registry.redhat.io/rhel9/rhel-guest-image:latest
```

```
    - name: registry.redhat.io/rhel8/rhel-guest-image:latest
```

```
    # Support and Troubleshooting Tools
```

```
    - name: registry.redhat.io/rhel8/support-tools:latest
```

```
    - name: registry.redhat.io/rhel9/support-tools:latest
```

```
    # Must-Gather Images for Debugging
```

```
    - name: registry.redhat.io/openshift4/ose-must-gather:latest
```

```
    - name: registry.redhat.io/container-native-virtualization/cnv-must-gather-
```

```
rhel9:v4.18
```

```
    - name: registry.redhat.io/odf4/odf-must-gather-rhel9:v4.18
```

```
    # NFS CSI Driver Dependencies
```

```
    - name: registry.k8s.io/sig-storage/csi-provisioner:v5.2.0
```

```
    - name: registry.k8s.io/sig-storage/csi-attacher:v4.8.0
```

```
    - name: registry.k8s.io/sig-storage/csi-resizer:v1.13.2
```

```
    - name: registry.k8s.io/sig-storage/csi-snapshotter:v8.2.0
```

```
    - name: registry.k8s.io/sig-storage/csi-node-driver-registrar:v2.13.0
```

```
operators:
```

Red Hat Operator Catalog

- catalog: registry.redhat.io/redhat/redhat-operator-index:v4.18

packages:

Core Platform Services

- name: advanced-cluster-management
- name: multicluster-engine
- name: multicluster-global-hub-operator-rh
- name: openshift-cert-manager-operator
- name: nfd
- name: windows-machine-config-operator

Storage and Data Management

- name: odf-operator
- name: odf-csi-addons-operator
- name: odf-dependencies
- name: odf-prometheus-operator
- name: odf-multicluster-orchestrator
- name: rook-ceph-operator
- name: redhat-oadp-operator

Virtualization Components

- name: kubevirt-hyperconverged
- name: kubernetes-nmstate-operator
- name: quay-operator
- name: quay-bridge-operator

Observability and Logging

- name: cluster-logging
- name: openshift-logging
- name: cluster-observability-operator
- name: loki-operator

CI/CD and Developer Tools

- name: openshift-pipelines-operator-rh
- name: openshift-gitops-operator
- name: web-terminal
- name: ansible-automation-platform-operator

Security and Compliance

- name: compliance-operator
- name: container-security-operator
- name: security-profiles-operator

Certified Operator Catalog

- catalog: registry.redhat.io/redhat/certified-operator-index:v4.18

packages:

Storage Solutions

```
- name: trident-operator
EOF
```

Step 1.3: Mirror Images to Disk

Purpose: Download all required images to local storage for transfer to air-gapped environment

```
# Authenticate with Red Hat registries
podman login registry.redhat.io
podman login quay.io

# Create mirror directory
mkdir -p /tmp/openshift-mirror

# Mirror images to disk using oc-mirror v2
echo "🔄 Starting image mirroring to disk (this may take several hours)..."
oc-mirror -c imageset-config.yaml file:///tmp/openshift-mirror --v2

# Verify mirror completed successfully
if [ $? -eq 0 ]; then
    echo "✅ Image mirroring completed successfully"
else
    echo "❌ Image mirroring failed - check logs and retry"
    exit 1
fi

# Create compressed archive for transfer
echo "📦 Creating transfer archive..."
cd /tmp/openshift-mirror
tar -czf openshift-4.18-fips-mirror.tar.gz working-dir/

# Verify archive created
ls -lh openshift-4.18-fips-mirror.tar.gz
echo "✅ Mirror archive ready for transfer: openshift-4.18-fips-mirror.tar.gz"
echo "📄 Archive size: $(du -h openshift-4.18-fips-mirror.tar.gz | cut -f1)"
```

Step 1.4: Transfer to Air-Gapped Environment

Purpose: Move the mirrored content to the disconnected environment

Transfer methods (choose based on your environment):

Option 1: Physical media transfer

```
echo "🌀 Copy archive to removable media for physical transfer"
cp openshift-4.18-fips-mirror.tar.gz /path/to/removable/media/
```

Option 2: Secure file transfer (if temporary connection allowed)

```
# scp openshift-4.18-fips-mirror.tar.gz user@octools.test.com:/tmp/
```

```
echo "⚠️ IMPORTANT: Transfer the following to air-gapped environment:"
```

```
echo "  - openshift-4.18-fips-mirror.tar.gz"
```

```
echo "  - imageset-config.yaml"
```

```
echo "  - pull-secret.json"
```

PHASE 2: AIR-GAPPED INFRASTRUCTURE PREPARATION

Context: This phase is performed INSIDE the air-gapped environment to prepare the base infrastructure that will support the OpenShift cluster.

Step 2.1: Prepare Bare Metal Infrastructure

Purpose: Wipe and prepare the physical servers that will host OpenShift

For each bare metal server (host01, host02, host03, nfs-server):

1. Boot from network or USB and wipe all disks

```
echo "🔧 Wiping bare metal servers..."
```

On each server, wipe the OS disk

```
sudo dd if=/dev/zero of=/dev/sda bs=1M count=1000
```

```
sudo wipefs -a /dev/sda
```

On OpenShift hosts, also wipe the data disk

```
sudo dd if=/dev/zero of=/dev/sdb bs=1M count=1000
```

```
sudo wipefs -a /dev/sdb
```

Configure static IP addressing for each server

```
# host01: 192.168.1.21
```

```
# host02: 192.168.1.22
```

```
# host03: 192.168.1.23
```

```
# nfs-server: 192.168.1.15
```

Step 2.2: Create vSphere VMs

Purpose: Create the supporting virtual machines needed for OpenShift deployment

Create VMs with the following specifications:

- Hostname: dns
 - CPU: 4vCPU
 - MEM: 4GB
 - DISK: 100GB
- Hostname: content
 - CPU: 4vCPU
 - MEM: 4GB
 - DISK: 100GB
 - OS: REHL9
- Hostname: mirror-registry
 - CPU: 8vCPU
 - MEM: 16GB
 - DISK: 1TB
 - OS: REHL9
- Hostname: octools
 - CPU: 8vCPU
 - MEM: 16GB
 - DISK: 200GB
 - OS: REHL9
- Hostname: bootstrap
 - CPU: 4vCPU
 - MEM: 16GB
 - DISK: 120GB
 - OS: RHCOS

PHASE 3: SUPPORTING SERVICES CONFIGURATION

Context: Configure the essential services (DNS, NTP, NFS, HTTP, Registry) that OpenShift requires to function in the disconnected environment.

Step 3.1: Configure DNS and NTP Server

Purpose: Provide name resolution and time synchronization for all OpenShift components

Install required packages

```
sudo dnf install -y bind bind-utils chrony
```

```
echo "🕒 Configuring internal NTP server..."
```

Configure chronyd as NTP server for the air-gapped environment

```
sudo tee /etc/chrony.conf << 'EOF'
```



```
# Use public servers as upstream (only during initial setup if internet available)
# Comment out these lines after initial sync in air-gapped environment
pool 2.rhel.pool.ntp.org iburst

# Allow clients from local network
allow 192.168.1.0/24

# Serve time even if not synchronized to a source (important for air-gapped)
local stratum 3

# Record the rate at which the system clock gains/loses time
driftfile /var/lib/chrony/drift

# Allow the system clock to be stepped in the first three updates
makestep 1.0 3

# Enable kernel synchronization of the real-time clock (RTC)
rtcsync

# Increase the minimum number of selectable sources required to adjust
# the system clock
minsources 2

# Allow NTP client access from local network
clientloglimit 1000000

# Specify directory for log files
logdir /var/log/chrony

# Select which information is logged
log measurements statistics tracking
EOF

# Enable and start chronyd
sudo systemctl enable --now chronyd

# Configure firewall for NTP
sudo firewall-cmd --permanent --add-service=ntp
sudo firewall-cmd --reload

echo "🌐 Configuring DNS server..."
# Configure BIND DNS server
sudo tee /etc/named.conf << 'EOF'
options {
    listen-on port 53 { 127.0.0.1; 192.168.1.13; };
    listen-on-v6 port 53 { ::1; };
```

```
directory "/var/named";
dump-file "/var/named/data/cache_dump.db";
statistics-file "/var/named/data/named_stats.txt";
memstatistics-file "/var/named/data/named_mem_stats.txt";
allow-query { localhost; 192.168.1.0/24; };
recursion yes;

forwarders {
    8.8.8.8;
    8.8.4.4;
};

dnssec-enable yes;
dnssec-validation yes;
};

logging {
    channel default_debug {
        file "data/named.run";
        severity dynamic;
    };
};

zone "." IN {
    type hint;
    file "named.ca";
};

include "/etc/named.rfc1912.zones";
include "/etc/named.root.key";

zone "test.com" IN {
    type master;
    file "test.com.zone";
    allow-update { none; };
};

zone "1.168.192.in-addr.arpa" IN {
    type master;
    file "1.168.192.in-addr.arpa.zone";
    allow-update { none; };
};
EOF

# Create forward DNS zone
sudo tee /var/named/test.com.zone << 'EOF'
$TTL 86400
```

```
@ IN SOA dns.test.com. admin.test.com. (
    2024010101 ; Serial
    3600       ; Refresh
    1800       ; Retry
    1209600    ; Expire
    86400 )    ; Minimum TTL

; Name servers
@ IN NS      dns.test.com.

; Supporting Infrastructure
dns      IN A    192.168.1.13
octools  IN A    192.168.1.11
tmpregistry IN A  192.168.1.10
content  IN A    192.168.1.13
nfs      IN A    192.168.1.15

; OpenShift Infrastructure
host01.ove IN A  192.168.1.21
host02.ove IN A  192.168.1.22
host03.ove IN A  192.168.1.23
bootstrap.ove IN A 192.168.1.25

; OpenShift API and Apps
api.ove IN A    192.168.1.30
api-int.ove IN A 192.168.1.30
*.apps.ove IN A  192.168.1.31

# Create reverse DNS zone
sudo tee /var/named/1.168.192.in-addr.arpa.zone << 'EOF'
$TTL 86400
@ IN SOA dns.test.com. admin.test.com. (
    2024010101 ; Serial
    3600       ; Refresh
    1800       ; Retry
    1209600    ; Expire
    86400 )    ; Minimum TTL

@ IN NS      dns.test.com.

; PTR records
13 IN PTR    dns.test.com.
11 IN PTR    octools.test.com.
10 IN PTR    tmpregistry.test.com.
13 IN PTR    content.test.com.
15 IN PTR    nfs.test.com.
21 IN PTR    host01.ove.test.com.
```

```
22      IN  PTR      host02.ove.test.com.
23      IN  PTR      host03.ove.test.com.
25      IN  PTR      bootstrap.ove.test.com.
30      IN  PTR      api.ove.test.com.
31      IN  PTR      apps.ove.test.com.
EOF
```

```
# Set proper permissions and start DNS
sudo chown named:named /var/named/*.zone
sudo chmod 644 /var/named/*.zone
```

```
# Validate DNS configuration
named-checkconf
named-checkzone test.com /var/named/test.com.zone
named-checkzone 1.168.192.in-addr.arpa /var/named/1.168.192.in-addr.arpa.zone
```

```
# Start and enable DNS services
sudo systemctl enable --now named
```

```
# Configure firewall
sudo firewall-cmd --permanent --add-service=dns
sudo firewall-cmd --reload
```

Step 3.2: Configure NFS Storage Server

Purpose: Provide persistent storage for OpenShift components and applications

```
# Install NFS packages on the dedicated bare metal server
sudo dnf install -y nfs-utils
```

```
echo "📁 Configuring NFS storage with LVM..."
# Create logical volumes for different storage tiers
sudo pvcreate /dev/sdb # Assuming 3TB disk is /dev/sdb
sudo vgcreate storage-vg /dev/sdb
```

```
# Create logical volumes for SSD and HDD storage classes
sudo lvcreate -L 1500G -n ssd-lv storage-vg
sudo lvcreate -L 1400G -n hdd-lv storage-vg
```

```
# Format filesystems
sudo mkfs.xfs /dev/storage-vg/ssd-lv
sudo mkfs.xfs /dev/storage-vg/hdd-lv
```

```
# Create mount points
sudo mkdir -p /exports/nfs-ssd
sudo mkdir -p /exports/nfs-hdd
```

```
# Add to fstab for persistent mounting
echo "/dev/storage-vg/ssd-lv /exports/nfs-ssd xfs defaults 0 0" | sudo tee -a /etc/fstab
echo "/dev/storage-vg/hdd-lv /exports/nfs-hdd xfs defaults 0 0" | sudo tee -a /etc/fstab

# Mount filesystems
sudo mount -a

# Configure NFS exports for OpenShift
sudo tee /etc/exports << 'EOF'
/exports/nfs-ssd 192.168.1.0/24(rw,sync,no_root_squash,no_all_squash,security_label)
/exports/nfs-hdd 192.168.1.0/24(rw,sync,no_root_squash,no_all_squash,security_label)
EOF

# Enable and start NFS services
sudo systemctl enable --now rpcbind
sudo systemctl enable --now nfs-server

# Export filesystems
sudo exportfs -rv

# Configure firewall for NFS
sudo firewall-cmd --permanent --add-service=nfs
sudo firewall-cmd --permanent --add-service=mountd
sudo firewall-cmd --permanent --add-service=rpc-bind
sudo firewall-cmd --reload

# Test NFS exports
showmount -e localhost
```

Step 3.3: Configure Content Server

Purpose: Host RHCOS images and ignition files for OpenShift node installation

```
# Install Apache HTTP Server
sudo dnf install -y httpd

echo "📁 Configuring content server directories..."
# Create content directories
sudo mkdir -p /var/www/html/rhcos
sudo mkdir -p /var/www/html/ignition
sudo mkdir -p /var/www/html/manifests

# Configure Apache
sudo systemctl enable --now httpd

# Configure firewall
sudo firewall-cmd --permanent --add-service=http
```

```
sudo firewall-cmd --reload
```

```
# Set SELinux context for web content
sudo setsebool -P httpd_read_user_content 1
sudo restorecon -R /var/www/html
```

Step 3.4: Configure Mirror Registry

Purpose: Host mirrored container images for OpenShift installation and operations

```
# Install Podman
sudo dnf install -y podman

echo "🏗️ Installing mirror registry..."
# Download and install mirror-registry (if not already transferred)
curl -LO https://console.redhat.com/openshift/downloads/tool-mirror-registry
tar xvf mirror-registry.tar.gz

# Install mirror registry with self-signed certificates
sudo ./mirror-registry install \
  --quayHostname tmpregistry.test.com \
  --quayRoot /opt/quay-install

# Note: Save the generated credentials output from installation
echo "⚠️ IMPORTANT: Save the generated registry credentials!"

# Configure firewall
sudo firewall-cmd --permanent --add-port=8443/tcp
sudo firewall-cmd --reload

# Test registry access
echo "🔑 Testing registry access..."
podman login -u init -p <generated_password> tmpregistry.test.com:8443 --tls-verify=false
```

PHASE 4: IMAGE MIRRORING AND TRANSFER

Context: Transfer the mirrored images from external environment and populate the internal mirror registry.

Step 4.1: Prepare FIPS-Enabled Tools Host

Purpose: Set up the host that will run OpenShift installation tools with FIPS compliance

⚠️ **CRITICAL:** This host MUST be FIPS-enabled before running any OpenShift installation commands.

```

echo "🔒 Enabling FIPS mode on octools host..."
# Enable FIPS mode FIRST - this requires a reboot
sudo fips-mode-setup --enable
echo "⚠️ REBOOTING to enable FIPS mode..."
sudo reboot

# After reboot, verify FIPS is enabled
echo "🔍 Verifying FIPS mode..."
fips-mode-setup --check
if [ $? -ne 0 ]; then
    echo "❌ ERROR: FIPS mode is not enabled!"
    exit 1
fi

# Verify FIPS enabled flag
cat /proc/sys/crypto/fips_enabled # Should show "1"

# Verify crypto policy is set to FIPS
update-crypto-policies --show # Should show "FIPS"

# Test FIPS compliance
echo "test" | openssl md5 2>&1 && echo "❌ ERROR: MD5 should be disabled in FIPS mode!"
|| echo "✅ MD5 properly disabled - FIPS compliant"

echo "📦 Installing OpenShift tools on FIPS-enabled host..."
# Install required packages
sudo dnf install -y curl wget tmux vim jq git

# Download and install OpenShift tools
# Download oc CLI
curl -L https://mirror.openshift.com/pub/openshift-v4/clients/ocp/stable-4.18/openshift-
client-linux.tar.gz -o oc.tar.gz
tar xvf oc.tar.gz
sudo mv oc kubectl /usr/local/bin/
chmod +x /usr/local/bin/oc /usr/local/bin/kubectl

# Download oc-mirror v2
curl -L https://console.redhat.com/openshift/downloads/tool-mirror-registry -o oc-
mirror.tar.gz
tar xvf oc-mirror.tar.gz
sudo mv oc-mirror /usr/local/bin/
chmod +x /usr/local/bin/oc-mirror

# Download openshift-install
curl -L https://mirror.openshift.com/pub/openshift-v4/clients/ocp/stable-4.18/openshift-
install-linux.tar.gz -o openshift-install.tar.gz

```

```
tar xvf openshift-install.tar.gz
sudo mv openshift-install /usr/local/bin/
chmod +x /usr/local/bin/openshift-install

# Download butane
curl -L https://mirror.openshift.com/pub/openshift-v4/clients/butane/latest/butane-amd64
-o butane
sudo mv butane /usr/local/bin/
chmod +x /usr/local/bin/butane

# Install Helm 3
curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash

# Configure system to use internal NTP
sudo tee /etc/chrony.conf << 'EOF'
# Use internal NTP server
server 192.168.1.13 iburst

# Record the rate at which the system clock gains/loses time
driftfile /var/lib/chrony/drift

# Allow the system clock to be stepped in the first three updates
makestep 1.0 3

# Enable kernel synchronization of the real-time clock (RTC)
rtcsync

# Specify directory for log files
logdir /var/log/chrony

# Select which information is logged
log measurements statistics tracking
EOF

sudo systemctl restart chronyd

# Verify time sync with internal NTP
chrony sources -v

# Verify tool installations
echo "🔧 Verifying tool installations..."
oc version
oc-mirror --help
openshift-install version
butane --version
helm version
```


Step 4.2: Transfer and Extract Mirror Content

Purpose: Move the mirrored images into the air-gapped environment and extract them

```
# Transfer the mirror archive to octools VM (via USB/network as appropriate)
echo "📦 Extracting mirror content..."
cd /tmp
# Assuming archive was transferred to /tmp/openshift-4.18-fips-mirror.tar.gz
tar -xzf openshift-4.18-fips-mirror.tar.gz

# Verify extraction
ls -la /tmp/working-dir/
```

Step 4.3: Configure Registry Authentication

Purpose: Set up authentication to push images to the internal mirror registry

```
echo "🔑 Configuring registry authentication..."
# Get registry certificate for trust
openssl s_client -connect tmpregistry.test.com:8443 -showcerts | awk '/BEGIN/,/END/{print $0}' | sudo tee /etc/pki/ca-trust/source/anchors/tmpregistry.crt
sudo update-ca-trust

# Configure authentication for internal registry
mkdir -p ~/.docker
cat << 'EOF' > ~/.docker/config.json
{
  "auths": {
    "tmpregistry.test.com:8443": {
      "auth": "<base64_encoded_init:password>",
      "email": "admin@test.com"
    }
  }
}
EOF

# Create authentication credentials (replace with actual base64 encoded credentials)
echo -n 'init:<registry_password>' | base64 -w0

# Test registry authentication
podman login tmpregistry.test.com:8443 --tls-verify=false
```

Step 4.4: Mirror Images to Internal Registry

Purpose: Push all OpenShift images from disk to the internal mirror registry

```
echo "🚀 Mirroring images to internal registry..."
# This process uploads all images to the internal registry
oc-mirror -c /tmp/imageset-config.yaml --from file:///tmp/working-dir
docker://tmpregistry.test.com:8443 --v2

# Verify mirroring completed successfully
if [ $? -eq 0 ]; then
    echo "✅ Image mirroring to internal registry completed successfully"
else
    echo "❌ Image mirroring failed - check logs and retry"
    exit 1
fi

# Verify cluster resources were generated
ls -la /tmp/working-dir/cluster-resources/
```

PHASE 5: OPENSHIFT INSTALLATION

Context: Install the OpenShift cluster using the mirrored images and FIPS-compliant configuration.

Step 5.1: Pre-Installation FIPS Validation

Purpose: Ensure FIPS compliance before starting OpenShift installation

```
echo "🔒 === MANDATORY FIPS Pre-Installation Verification ==="
echo "1. FIPS Mode Status:"
fips-mode-setup --check
if [ $? -ne 0 ]; then
    echo "❌ ERROR: FIPS mode is not enabled on this host!"
    echo "You MUST enable FIPS mode and reboot before proceeding:"
    echo "sudo fips-mode-setup --enable"
    echo "sudo reboot"
    exit 1
fi

echo "2. FIPS Enabled Flag:"
FIPS_ENABLED=$(cat /proc/sys/crypto/fips_enabled 2>/dev/null)
if [ "$FIPS_ENABLED" != "1" ]; then
    echo "❌ ERROR: FIPS is not properly enabled (fips_enabled = $FIPS_ENABLED)"
    exit 1
fi
```

```

echo "3. Crypto Policy:"
CRYPTO_POLICY=$(update-crypto-policies --show)
if [ "$CRYPTO_POLICY" != "FIPS" ]; then
    echo "❌ ERROR: Crypto policy is not set to FIPS (current: $CRYPTO_POLICY)"
    echo "Run: sudo update-crypto-policies --set FIPS"
    exit 1
fi

echo "4. OpenSSL FIPS Test:"
echo "test" | openssl md5 2>&1 && echo "❌ ERROR: MD5 should be disabled in FIPS mode!"
&& exit 1
echo "✅ MD5 properly disabled - FIPS compliant"

echo ""
echo "✅ FIPS verification PASSED - proceeding with OpenShift installation"
echo ""

```

Step 5.2: Download and Prepare RHCOS Images

Purpose: Obtain the Red Hat CoreOS images needed for node installation

```

echo "📁 Downloading RHCOS images..."
mkdir -p ~/openshift-install
cd ~/openshift-install

# Download RHCOS images for OpenShift 4.18
curl -L https://mirror.openshift.com/pub/openshift-
v4/dependencies/rhcos/4.18/latest/rhcos-live.x86_64.iso -o rhcos-live.iso
curl -L https://mirror.openshift.com/pub/openshift-
v4/dependencies/rhcos/4.18/latest/rhcos-metal.x86_64.raw.gz -o rhcos-metal.raw.gz

# Copy RHCOS images to content server for network access
scp rhcos-* root@content.test.com:/var/www/html/rhcos/

echo "✅ RHCOS images available on content server"

```

Step 5.3: Create Installation Configuration

Purpose: Generate the OpenShift installation configuration with FIPS enabled

```

echo "⚙️ Creating FIPS-enabled install-config.yaml..."
# Create install-config.yaml with FIPS enabled
cat << 'EOF' > install-config.yaml
apiVersion: v1
baseDomain: test.com
metadata:
  name: ove

```

```
networking:
  networkType: OVNKubernetes
  clusterNetwork:
    - cidr: 10.128.0.0/14
      hostPrefix: 23
  serviceNetwork:
    - 172.30.0.0/16
  machineNetwork:
    - cidr: 192.168.1.0/24
compute:
  - name: worker
    replicas: 0
controlPlane:
  name: master
  replicas: 3
  platform:
    baremetal: {}
platform:
  baremetal:
    apiVIPs:
      - 192.168.1.30
    ingressVIPs:
      - 192.168.1.31
    hosts:
      - name: host01
        role: master
        bmc:
          address: redfish://192.168.1.21
          username: admin
          password: password
          bootMACAddress: "aa:bb:cc:dd:ee:01"
          rootDeviceHints:
            deviceName: "/dev/sda"
      - name: host02
        role: master
        bmc:
          address: redfish://192.168.1.22
          username: admin
          password: password
          bootMACAddress: "aa:bb:cc:dd:ee:02"
          rootDeviceHints:
            deviceName: "/dev/sda"
      - name: host03
        role: master
        bmc:
          address: redfish://192.168.1.23
          username: admin
```

```

    password: password
    bootMACAddress: "aa:bb:cc:dd:ee:03"
    rootDeviceHints:
      deviceName: "/dev/sda"
fips: true
pullSecret: |
  <your_pull_secret_with_internal_registry>
sshKey: |
  <your_ssh_public_key>
imageContentSources:
- mirrors:
  - tmpregistry.test.com:8443/openshift/release-images
  source: quay.io/openshift-release-dev/ocp-release
- mirrors:
  - tmpregistry.test.com:8443/openshift/release
  source: quay.io/openshift-release-dev/ocp-v4.0-art-dev
additionalTrustBundle: |
  -----BEGIN CERTIFICATE-----
  <registry_certificate_content>
  -----END CERTIFICATE-----
EOF

```

```

# Verify FIPS is enabled in configuration
grep -q "fips: true" install-config.yaml && echo "✅ FIPS enabled in install-config.yaml"
|| echo "❌ ERROR: FIPS not enabled in install-config.yaml"

```

Step 5.4: Generate Custom Manifests

Purpose: Create custom configurations for internal NTP and FIPS compliance

```

echo "📄 Creating cluster manifests..."
openshift-install create manifests

```

```

echo "🕒 Creating custom NTP configuration manifests..."
# Create custom NTP configuration for masters
cat << 'EOF' > manifests/99-master-chrony-configuration.yaml
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: master
  name: 99-master-chrony-configuration
spec:
  config:
    ignition:
      version: 3.2.0
    storage:

```

```

files:
- contents:
    source: data:text/plain;charset=utf-
8;base64,IyBVc2Ugaw50ZXJuYWwgTlRQIHN1cnZlcgpzZXJ2ZXIgmTkyLjE2OC4xLjEzIGlidXJzdAoKIyBSZWNV
cmQgdGhlIHJhdGUgYXQgd2hpY2ggdGhlIHN5c3RlbSBjbG9jayBnYWlucy9sb3NlcyB0aW1lcmRyaWZ0ZmlsZSAvd
mFyL2xpYi9jaHJvbnkvZHJpZnQKCiMgQWxs3cgdGhlIHN5c3RlbSBjbG9jayB0byBiZSBzdGVwcGVkIGluIHRoZS
BmaXJzdCB0aHJlZSB1cGRhdGVzCm1ha2VzdGVwIDEuMCAzCgojIEVuYWJsZSBzZXJ2ZWwg3luY2hyb25pemF0aW9
uIG9mIHRoZSBzYWFsLXRpbWUgY2xvY2sgKFJUQykKcnRjc3luYwoKIyBjbG9jayB0byBiZSBzdGVwcGVkIGluIHRoZS
ZXIgb2Ygc2VsZW90YWJsZSBzb3VyY2VzIHJlcXVpcmkvIHRvIGFkanVzdAoJIHRoZSBzeXN0ZW0gY2xvY2sKbWluc
291cmNlcyAxGgojIFNwZWNPZnkgZGlyZW90b3J5IGZvciBsb2cgZmlsZXMKbG9nZGlyIC92YXlvcG9nL2Nocm9ueQ
oKIyBTZw1Y3Qgd2hpY2ggaw5mb3JtYXRpb24gaXMgbG9nZ2VkCmxvZyBtZWZzdXJlbWVudHMgc3RhdGlzdGljcyB
0cmFja2luZwo=
    mode: 420
    overwrite: true
    path: /etc/chrony.conf
systemd:
  units:
    - enabled: true
      name: chronyd.service
EOF

# Create custom NTP configuration for workers
cat << 'EOF' > manifests/99-worker-chrony-configuration.yaml
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: worker
  name: 99-worker-chrony-configuration
spec:
  config:
    ignition:
      version: 3.2.0
    storage:
      files:
        - contents:
            source: data:text/plain;charset=utf-
8;base64,IyBVc2Ugaw50ZXJuYWwgTlRQIHN1cnZlcgpzZXJ2ZXIgmTkyLjE2OC4xLjEzIGlidXJzdAoKIyBSZWNV
cmQgdGhlIHJhdGUgYXQgd2hpY2ggdGhlIHN5c3RlbSBjbG9jayBnYWlucy9sb3NlcyB0aW1lcmRyaWZ0ZmlsZSAvd
mFyL2xpYi9jaHJvbnkvZHJpZnQKCiMgQWxs3cgdGhlIHN5c3RlbSBjbG9jayB0byBiZSBzdGVwcGVkIGluIHRoZS
BmaXJzdCB0aHJlZSB1cGRhdGVzCm1ha2VzdGVwIDEuMCAzCgojIEVuYWJsZSBzZXJ2ZWwg3luY2hyb25pemF0aW9
uIG9mIHRoZSBzYWFsLXRpbWUgY2xvY2sgKFJUQykKcnRjc3luYwoKIyBjbG9jayB0byBiZSBzdGVwcGVkIGluIHRoZS
ZXIgb2Ygc2VsZW90YWJsZSBzb3VyY2VzIHJlcXVpcmkvIHRvIGFkanVzdAoJIHRoZSBzeXN0ZW0gY2xvY2sKbWluc
291cmNlcyAxGgojIFNwZWNPZnkgZGlyZW90b3J5IGZvciBsb2cgZmlsZXMKbG9nZGlyIC92YXlvcG9nL2Nocm9ueQ
oKIyBTZw1Y3Qgd2hpY2ggaw5mb3JtYXRpb24gaXMgbG9nZ2VkCmxvZyBtZWZzdXJlbWVudHMgc3RhdGlzdGljcyB
0cmFja2luZwo=
            mode: 420

```

```

        overwrite: true
        path: /etc/chrony.conf
systemd:
  units:
    - enabled: true
      name: chronyd.service
EOF

```

```

echo "🔒 Creating FIPS kernel argument manifests..."
# Create FIPS compliance manifest for additional kernel arguments
cat << 'EOF' > manifests/99-master-fips-kargs.yaml
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: master
    name: 99-master-fips-kargs
spec:
  kernelArguments:
    - fips=1
EOF

```

```

cat << 'EOF' > manifests/99-worker-fips-kargs.yaml
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: worker
    name: 99-worker-fips-kargs
spec:
  kernelArguments:
    - fips=1
EOF

```

Step 5.5: Generate Ignition Configurations

Purpose: Create the ignition files that will configure the OpenShift nodes

```

echo "🔥 Generating ignition configurations..."
# Create ignition configs
openshift-install create ignition-configs

# Verify ignition files were created
ls -la *.ign
echo "✅ Ignition files generated:"
echo "  - bootstrap.ign (for bootstrap node)"
echo "  - master.ign (for master nodes)"

```

```
echo "    - worker.ign (for worker nodes)"

# Copy ignition files to content server for network access during installation
echo "📁 Copying ignition files to content server..."
scp *.ign root@content.test.com:/var/www/html/ignition/
```

Step 5.6: Install Bootstrap Node

Purpose: Create and start the temporary bootstrap node that orchestrates cluster installation

```
echo "🚀 Starting bootstrap node installation..."
echo "📁 Bootstrap installation steps:"
echo "1. Create bootstrap VM with 4c, 16GB, 120GB SSD"
echo "2. Boot from RHCOS live ISO"
echo "3. Run coreos-installer with bootstrap ignition"

# Manual steps to perform on bootstrap VM:
echo "
🖥 On the bootstrap VM console:
sudo coreos-installer install /dev/sda \\\
  --ignition-url http://content.test.com/ignition/bootstrap.ign \\\
  --insecure-ignition

# Wait for installation to complete, then reboot
sudo reboot
"

echo "⌚ Waiting for bootstrap node to become ready..."
echo "    Monitor logs: journalctl -b -f -u release-image.service -u bootkube.service"
```

Step 5.7: Install Control Plane Nodes

Purpose: Install the three master nodes that will form the OpenShift control plane

```
echo "🎯 Installing master nodes..."
echo "📁 Master node installation steps:"

# For each master node (host01, host02, host03):
echo "
🖥 For each master node, boot from RHCOS live ISO and run:

# On host01 (192.168.1.21):
sudo coreos-installer install /dev/sda \\\
  --ignition-url http://content.test.com/ignition/master.ign \\\
  --insecure-ignition
```



```
# On host02 (192.168.1.22):
sudo coreos-installer install /dev/sda \
  --ignition-url http://content.test.com/ignition/master.ign \
  --insecure-ignition

# On host03 (192.168.1.23):
sudo coreos-installer install /dev/sda \
  --ignition-url http://content.test.com/ignition/master.ign \
  --insecure-ignition

# After installation, reboot each node
sudo reboot
"

echo "🕒 Waiting for master nodes to join the cluster..."
```

Step 5.8: Monitor Installation Progress

Purpose: Track the installation progress and handle the bootstrap removal

```
echo "🇺🇸 Monitoring OpenShift installation progress..."

# Monitor bootstrap completion
echo "🕒 Waiting for bootstrap to complete..."
openshift-install wait-for bootstrap-complete --log-level=info

if [ $? -eq 0 ]; then
  echo "✅ Bootstrap phase completed successfully"
  echo "🗑️ Bootstrap node can now be safely removed"
else
  echo "❌ Bootstrap phase failed - check logs"
  exit 1
fi

# At this point, remove or shutdown the bootstrap VM
echo "🗑️ Remove bootstrap VM from load balancer and power down"

# Wait for installation to complete
echo "🕒 Waiting for installation to complete..."
openshift-install wait-for install-complete --log-level=info

if [ $? -eq 0 ]; then
  echo "🎉 OpenShift installation completed successfully!"
  echo "🗑️ Cluster access information:"
  cat auth/kubeconfig
else
  echo "❌ Installation failed - check logs and troubleshoot"
```

```
    exit 1
fi

echo "✅ OpenShift 4.18 cluster with FIPS compliance is now running"
```

PHASE 6: POST-INSTALLATION CONFIGURATION

Context: Configure the newly installed OpenShift cluster with mirrored content and basic settings.

Step 6.1: Configure Cluster Authentication

Purpose: Set up access to the newly installed cluster

```
echo "🔑 Configuring cluster access..."
# Set up kubeconfig for cluster access
export KUBECONFIG=~/.openshift-install/auth/kubeconfig

# Test cluster access
oc get nodes
oc get clusterversion

# Show cluster information
echo "📋 Cluster Information:"
oc cluster-info
oc get clusteroperators

echo "✅ Cluster access configured"
```

Step 6.2: Apply Mirrored Content Configuration

Purpose: Configure the cluster to use the mirrored images and operator catalogs

```
echo "🔄 Applying mirrored content configuration..."
# Apply the generated cluster resources from oc-mirror
oc apply -f /tmp/working-dir/cluster-resources/

# Verify ImageDigestMirrorSet and ImageTagMirrorSet are applied
oc get imagedigestmirrorset
oc get imagetagmirrorset

# Verify CatalogSources are available
oc get catalogsource -n openshift-marketplace

# Wait for catalog sources to be ready
echo "⌚ Waiting for catalog sources to become ready..."
```

```
oc wait --for=condition=Ready catalogsource --all -n openshift-marketplace --timeout=300s  
echo "✅ Mirrored content configuration applied successfully"
```

Step 6.3: Validate FIPS Compliance

Purpose: Ensure the cluster is running in FIPS mode as expected

```
echo "🔒 Validating cluster FIPS compliance..."  
  
# Check FIPS status on all nodes  
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}'); do  
    echo "🔍 Checking FIPS on $node:"  
    oc debug node/$node -- chroot /host fips-mode-setup --check  
    echo "FIPS kernel arg: $(oc debug node/$node -- chroot /host cat /proc/cmdline  
2>/dev/null | grep -o 'fips=[0-9]*')"  
    echo "FIPS enabled flag: $(oc debug node/$node -- chroot /host cat  
/proc/sys/crypto/fips_enabled 2>/dev/null)"  
    echo "---"  
done  
  
# Verify cluster-wide FIPS configuration  
oc get proxy cluster -o yaml | grep -i fips  
  
echo "✅ FIPS compliance validation completed"
```

PHASE 7: PLATFORM PLUS COMPONENTS INSTALLATION

Context: Install the additional Platform Plus components including storage, virtualization, and management tools.

Step 7.1: Configure NFS Storage Classes

Purpose: Set up persistent storage using the NFS server for OpenShift components

```
echo "📁 Configuring NFS storage classes..."

# Deploy NFS CSI driver
echo "📦 Installing NFS CSI driver..."
oc create namespace nfs-csi-driver

# Create NFS CSI driver deployment
cat << 'EOF' | oc apply -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: csi-nfs-controller
  namespace: nfs-csi-driver
spec:
  replicas: 1
  selector:
    matchLabels:
      app: csi-nfs-controller
  template:
    metadata:
      labels:
        app: csi-nfs-controller
    spec:
      serviceAccountName: csi-nfs-controller-sa
      containers:
        - name: csi-provisioner
          image: tmpregistry.test.com:8443/sig-storage/csi-provisioner:v5.2.0
          args:
            - "--csi-address=$(ADDRESS)"
            - "--v=2"
            - "--leader-election=true"
            - "--leader-election-namespace=nfs-csi-driver"
          env:
            - name: ADDRESS
              value: /csi/csi.sock
          volumeMounts:
            - mountPath: /csi
              name: socket-dir
        - name: csi-nfs-driver
          image: tmpregistry.test.com:8443/sig-storage/nfsplugin:v4.9.0
          args:
            - "--endpoint=$(CSI_ENDPOINT)"
            - "--nodeid=$(NODE_ID)"
```

```
- "--v=2"
env:
- name: CSI_ENDPOINT
  value: unix:///csi/csi.sock
- name: NODE_ID
  valueFrom:
    fieldRef:
      fieldPath: spec.nodeName
volumeMounts:
- mountPath: /csi
  name: socket-dir
volumes:
- name: socket-dir
  emptyDir: {}
```

EOF

```
# Create SSD storage class
echo "⚡ Creating NFS SSD storage class..."
cat << 'EOF' | oc apply -f -
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: nfs-ssd
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
provisioner: nfs.csi.k8s.io
parameters:
  server: 192.168.1.15
  share: /exports/nfs-ssd
  mountPermissions: "0755"
reclaimPolicy: Delete
volumeBindingMode: Immediate
mountOptions:
- hard
- nfsvers=4.1
EOF
```

```
# Create HDD storage class
echo "💿 Creating NFS HDD storage class..."
cat << 'EOF' | oc apply -f -
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: nfs-hdd
provisioner: nfs.csi.k8s.io
parameters:
```

```
server: 192.168.1.15
share: /exports/nfs-hdd
mountPermissions: "0755"
reclaimPolicy: Delete
volumeBindingMode: Immediate
mountOptions:
  - hard
  - nfsvers=4.1
EOF
```

```
echo "✅ NFS storage classes configured"
```

Step 7.2: Configure Image Registry with Persistent Storage

Purpose: Set up the internal image registry to use persistent storage

```
echo "🖼️ Configuring image registry with persistent storage..."
```

```
# Create PVC for image registry
cat << 'EOF' | oc apply -f -
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: image-registry-storage
  namespace: openshift-image-registry
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 500Gi
  storageClassName: nfs-ssd
EOF
```

```
# Configure image registry to use NFS storage
oc patch configs.imageregistry.operator.openshift.io cluster --type merge --patch
'{"spec":{"storage":{"pvc":{"claim":"image-registry-storage"}}}}'
```

```
# Set image registry to managed
oc patch configs.imageregistry.operator.openshift.io cluster --type merge --patch
'{"spec":{"managementState":"Managed"}}'
```

```
# Verify image registry deployment
echo "⌚ Waiting for image registry to become ready..."
oc wait --for=condition=Available deployment/image-registry -n openshift-image-registry -
-timeout=300s
```

```
echo "✅ Image registry configured with persistent storage"
```

Step 7.3: Install OpenShift Virtualization

Purpose: Enable VM workloads alongside containers

```
echo "🖥 Installing OpenShift Virtualization..."
```

```
# Create namespace for OpenShift Virtualization
oc create namespace openshift-cnv
```

```
# Create OperatorGroup
cat << 'EOF' | oc apply -f -
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: kubevirt-hyperconverged-group
  namespace: openshift-cnv
spec:
  targetNamespaces:
    - openshift-cnv
EOF
```

```
# Create Subscription for OpenShift Virtualization
echo "📦 Installing kubevirt-hyperconverged operator..."
cat << 'EOF' | oc apply -f -
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: hco-operatorhub
  namespace: openshift-cnv
spec:
  source: redhat-operators
  sourceNamespace: openshift-marketplace
  name: kubevirt-hyperconverged
  channel: "stable"
EOF
```

```
# Wait for operator installation
echo "⌚ Waiting for OpenShift Virtualization operator to install..."
oc wait --for=condition=Succeeded csv -l operators.coreos.com/kubevirt-
hyperconverged.openshift-cnv -n openshift-cnv --timeout=600s
```

```
# Create HyperConverged instance with FIPS-compatible settings
echo "🔧 Creating HyperConverged instance..."
```

```
cat << 'EOF' | oc apply -f -
apiVersion: hco.kubevirt.io/v1beta1
kind: HyperConverged
metadata:
  name: kubevirt-hyperconverged
  namespace: openshift-cnv
spec:
  infra:
    nodePlacement:
      nodeSelector:
        node-role.kubernetes.io/worker: ""
  workloads:
    nodePlacement:
      nodeSelector:
        node-role.kubernetes.io/worker: ""
  featureGates:
    deployTektonTaskResources: false
    disableMDevConfiguration: false
    enableCommonBootImageImport: true
    deployVmConsoleProxy: true
  certConfig:
    ca:
      duration: 8760h0m0s
      renewBefore: 720h0m0s
    server:
      duration: 8760h0m0s
      renewBefore: 720h0m0s
  tlsSecurityProfile:
    type: Intermediate
    intermediate:
      ciphers:
        - ECDHE-ECDSA-AES128-GCM-SHA256
        - ECDHE-RSA-AES128-GCM-SHA256
        - ECDHE-ECDSA-AES256-GCM-SHA384
        - ECDHE-RSA-AES256-GCM-SHA384
        - ECDHE-ECDSA-CHACHA20-POLY1305
        - ECDHE-RSA-CHACHA20-POLY1305
      minTLSVersion: VersionTLS12
EOF

# Monitor deployment progress
echo "⌚ Waiting for OpenShift Virtualization to become ready..."
oc wait --for=condition=Available hco kubevirt-hyperconverged -n openshift-cnv --
timeout=1200s

echo "✅ OpenShift Virtualization installed and ready"
```


Step 7.4: Install Advanced Cluster Management

Purpose: Enable multi-cluster management capabilities

```
echo "🌐 Installing Advanced Cluster Management..."

# Create namespace
oc create namespace open-cluster-management

# Create OperatorGroup
cat << 'EOF' | oc apply -f -
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: open-cluster-management
  namespace: open-cluster-management
spec:
  targetNamespaces:
    - open-cluster-management
EOF

# Install ACM operator
echo "📦 Installing advanced-cluster-management operator..."
cat << 'EOF' | oc apply -f -
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: advanced-cluster-management
  namespace: open-cluster-management
spec:
  source: redhat-operators
  sourceNamespace: openshift-marketplace
  name: advanced-cluster-management
  channel: "release-2.12"
EOF

# Wait for operator installation
echo "⌚ Waiting for ACM operator to install..."
oc wait --for=condition=Succeeded csv -l operators.coreos.com/advanced-cluster-
management.open-cluster-management -n open-cluster-management --timeout=600s
```

Step 7.5: Configure Monitoring with Persistent Storage

Purpose: Set up monitoring stack with persistent storage for metrics retention

```
echo "🇮🇹 Configuring monitoring with persistent storage..."
```

```
# Create monitoring storage configuration
cat << 'EOF' | oc apply -f -
apiVersion: v1
kind: ConfigMap
metadata:
  name: cluster-monitoring-config
  namespace: openshift-monitoring
data:
  config.yaml: |
    prometheusK8s:
      volumeClaimTemplate:
        spec:
          storageClassName: nfs-ssd
          resources:
            requests:
              storage: 100Gi
    alertmanagerMain:
      volumeClaimTemplate:
        spec:
          storageClassName: nfs-ssd
          resources:
            requests:
              storage: 10Gi
    thanosRuler:
      volumeClaimTemplate:
        spec:
          storageClassName: nfs-ssd
          resources:
            requests:
              storage: 10Gi
EOF

# Create user workload monitoring configuration
cat << 'EOF' | oc apply -f -
apiVersion: v1
kind: ConfigMap
metadata:
  name: user-workload-monitoring-config
  namespace: openshift-user-workload-monitoring
data:
  config.yaml: |
    prometheus:
      volumeClaimTemplate:
        spec:
          storageClassName: nfs-ssd
          resources:
```

```

      requests:
        storage: 50Gi
    thanosRuler:
      volumeClaimTemplate:
        spec:
          storageClassName: nfs-ssd
          resources:
            requests:
              storage: 10Gi

```

EOF

echo "✅ Monitoring configured with persistent storage"

PHASE 8: VALIDATION AND TESTING

Context: Validate that all components are working correctly and test FIPS compliance.

Step 8.1: Comprehensive FIPS Validation

Purpose: Ensure the entire cluster is FIPS-compliant

echo "🔒 === Comprehensive FIPS Compliance Validation ==="

```

# Create FIPS validation job
cat << 'EOF' | oc apply -f -
apiVersion: batch/v1
kind: Job
metadata:
  name: fips-validation
spec:
  template:
    spec:
      containers:
      - name: fips-checker
        image: tmpregistry.test.com:8443/ubi9/ubi:latest
        command:
        - /bin/bash
        - -c
        - |
          echo "=== FIPS Mode Validation ==="
          fips-mode-setup --check
          echo "Exit code: $?"
          echo ""
          echo "=== Kernel Command Line ==="
          cat /proc/cmdline
          echo ""

```

```

    echo "=== OpenSSL FIPS Test ==="
    openssl md5 /dev/null 2>&1 || echo "FIPS mode detected - MD5 disabled as
expected"
    echo ""
    echo "=== Crypto Policy ==="
    update-crypto-policies --show
    echo ""
    echo "=== FIPS Module Status ==="
    if [ -f /proc/sys/crypto/fips_enabled ]; then
        echo "FIPS enabled: $(cat /proc/sys/crypto/fips_enabled)"
    else
        echo "FIPS status file not found"
    fi
    securityContext:
        privileged: true
    restartPolicy: Never
    hostNetwork: true
    hostPID: true
    backoffLimit: 1
EOF

# Check FIPS validation results
echo "🕒 Waiting for FIPS validation to complete..."
oc wait --for=condition=Complete job/fips-validation --timeout=300s
oc logs job/fips-validation

# Clean up validation job
oc delete job fips-validation

# Verify FIPS compliance across all nodes
echo "🔍 Validating FIPS on all cluster nodes..."
for node in $(oc get nodes -o name); do
    echo "=== Checking FIPS on $node ==="
    oc debug $node -- chroot /host fips-mode-setup --check
done

echo "✅ FIPS validation completed"

```

Step 8.2: Test OpenShift Virtualization

Purpose: Verify VM capabilities are working correctly

```
echo "🖥️ Testing OpenShift Virtualization..."
```

```
# Create a test VM with FIPS compliance
```

```
cat << 'EOF' | oc apply -f -
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: test-vm-fips
  namespace: default
spec:
  running: true
  template:
    metadata:
      labels:
        kubevirt.io/vm: test-vm-fips
    spec:
      domain:
        cpu:
          cores: 2
        devices:
          disks:
            - disk:
                bus: virtio
                name: containerdisk
            - disk:
                bus: virtio
                name: cloudinitdisk
          interfaces:
            - name: default
              masquerade: {}
        machine:
          type: pc-q35-rhel9.2.0
        resources:
          requests:
            memory: 2Gi
      networks:
        - name: default
          pod: {}
      volumes:
        - containerDisk:
            image: tmpregistry.test.com:8443/rhel9/rhel-guest-image:latest
            name: containerdisk
        - cloudInitNoCloud:
            userData: |
              #cloud-config
            users:
              - name: rhel
                sudo: ['ALL=(ALL) NOPASSWD:ALL']
                ssh_authorized_keys:
                  - ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQACAQ...
```

```

        runcmd:
          - fips-mode-setup --enable
          - systemctl reboot
      name: cloudinitdisk
EOF

# Wait for VM to be ready
echo "🕒 Waiting for test VM to start..."
oc wait --for=condition=Ready vmi/test-vm-fips --timeout=300s

# Check VM status
oc get vm test-vm-fips
oc get vmi test-vm-fips

echo "✅ OpenShift Virtualization test VM created successfully"

```

Step 8.3: Test Storage Classes

Purpose: Verify persistent storage is working correctly

```

echo "💾 Testing storage classes..."

# Create test PVC for SSD storage
cat << 'EOF' | oc apply -f -
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: test-pvc-ssd-fips
  annotations:
    security.openshift.io/fips-compliant: "true"
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: nfs-ssd
EOF

# Create test PVC for HDD storage
cat << 'EOF' | oc apply -f -
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: test-pvc-hdd-fips
  annotations:

```

```

    security.openshift.io/fips-compliant: "true"
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: nfs-hdd
EOF

```

```

# Wait for PVCs to be bound
echo "🕒 Waiting for test PVCs to be bound..."
oc wait --for=condition=Bound pvc/test-pvc-ssd-fips --timeout=300s
oc wait --for=condition=Bound pvc/test-pvc-hdd-fips --timeout=300s

# Check PVC status
oc get pvc test-pvc-ssd-fips test-pvc-hdd-fips

echo "✅ Storage classes tested successfully"

```

Step 8.4: Test NTP Synchronization

Purpose: Verify all nodes are synchronized with internal NTP

```

echo "🕒 Testing NTP synchronization..."

# Create NTP test pod
cat << 'EOF' | oc apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: ntp-sync-test
spec:
  containers:
    - name: ntp-test
      image: tmpregistry.test.com:8443/ubi9/ubi:latest
      command:
        - /bin/bash
        - -c
        - |
          echo "=== NTP Synchronization Test ==="
          dnf install -y chrony
          echo "server 192.168.1.13 iburst" > /etc/chrony.conf
          systemctl start chronyd
          sleep 10
          echo "1. Chrony sources:"
          chrony sources -v

```

```

echo ""
echo "2. Chrony tracking:"
chrony tracking
echo ""
echo "3. Time status:"
timedatectl status
echo ""
echo "4. Testing NTP connectivity:"
chrony sourcestats
sleep 60
securityContext:
  privileged: true
hostNetwork: true
restartPolicy: Never
EOF

# Monitor NTP test
echo "🕒 Running NTP synchronization test..."
oc wait --for=condition=Succeeded pod/ntp-sync-test --timeout=300s
oc logs ntp-sync-test

# Clean up test pod
oc delete pod ntp-sync-test

echo "✅ NTP synchronization test completed"

```

Step 8.5: Validate Operator Installation

Purpose: Ensure all Platform Plus operators are functioning correctly

```

echo "🔧 Validating operator installations..."

# Check all operator statuses
echo "📋 Checking subscription statuses:"
oc get subscription -A

echo ""
echo "📋 Checking CSV statuses:"
oc get csv -A

echo ""
echo "📋 Checking operator pod statuses:"
oc get pods -n openshift-cnv | grep -E "(Running|Completed)"
oc get pods -n open-cluster-management | grep -E "(Running|Completed)"
echo "✅ Operator validation completed"

```


Step 8.6: Final Cluster Health Check

Purpose: Comprehensive cluster health and functionality validation

```
echo "🔧 Performing final cluster health check..."

# Check cluster operators
echo "🔧 Cluster Operators Status:"
oc get clusteroperators

# Check node status
echo "🖥️ Node Status:"
oc get nodes -o wide

# Check critical namespaces
echo "📦 Critical Namespace Status:"
for ns in openshift-kube-apiserver openshift-etcd openshift-kube-controller-manager
openshift-kube-scheduler; do
    echo "Namespace: $ns"
    oc get pods -n $ns | grep -v Completed
    echo "---"
done

# Check storage
echo "📁 Storage Status:"
oc get storageclass
oc get pv

# Check catalog sources
echo "📖 Catalog Sources:"
oc get catalogsouce -n openshift-marketplace

# Check virtualization
echo "🖥️ Virtualization Status:"
oc get hco -n openshift-cnv
oc get kubvirt -n openshift-cnv

echo "✅ Final cluster health check completed"
```

TROUBLESHOOTING GUIDE

Context: Common issues and their solutions for the FIPS-enabled disconnected OpenShift deployment.

DNS and Network Issues

```
echo "🌐 DNS Troubleshooting Commands:"
```

```
# Test DNS resolution from nodes
oc debug node/host01.ove.test.com -- chroot /host nslookup api.ove.test.com

# Check DNS configuration on nodes
oc debug node/host01.ove.test.com -- chroot /host cat /etc/resolv.conf

# Verify DNS server is responding
dig @192.168.1.13 api.ove.test.com

# Test etcd SRV records
dig @192.168.1.13 _etcd-server-ssl._tcp.ove.test.com SRV

# Check network connectivity between nodes
oc debug node/host01.ove.test.com -- chroot /host ping host02.ove.test.com
```

Registry and Image Issues

```
echo "🚧 Registry Troubleshooting Commands:"

# Check if registry certificate is trusted
oc debug node/host01.ove.test.com -- chroot /host openssl s_client -connect
tmpregistry.test.com:8443 -verify_return_error

# Check ImageContentSourcePolicy configuration
oc get imagecontentsourcepolicy
oc get imagedigestmirrorset
oc get imagetagmirrorset

# Verify mirror configuration
oc get images.config.openshift.io cluster -o yaml

# Check node's registries.conf
oc debug node/host01.ove.test.com -- chroot /host cat /etc/containers/registries.conf

# Test image pull from mirror registry
oc debug node/host01.ove.test.com -- chroot /host podman pull
tmpregistry.test.com:8443/ubi9/ubi:latest

# Check registry pod logs
oc logs -n openshift-image-registry deployment/image-registry
```

FIPS Compliance Issues

```
echo "🔒 FIPS Troubleshooting Commands:"

# Check FIPS mode status on nodes
oc debug node/host01.ove.test.com -- chroot /host fips-mode-setup --check
```

```
# Verify FIPS kernel arguments
oc debug node/host01.ove.test.com -- chroot /host cat /proc/cmdline | grep fips

# Check FIPS compliance for OpenSSL
oc debug node/host01.ove.test.com -- chroot /host openssl md5 /dev/null

# Verify FIPS mode in containers
oc run fips-test --image=tmpregistry.test.com:8443/ubi9/ubi:latest --rm -it --
restart=Never -- fips-mode-setup --check

# Check FIPS enabled services
oc debug node/host01.ove.test.com -- chroot /host systemctl status fips-mode-setup

# Verify crypto policies
oc debug node/host01.ove.test.com -- chroot /host update-crypto-policies --show

# Check for FIPS-related machine configs
oc get machineconfig | grep fips
```

NTP Synchronization Issues

```
echo "🕒 NTP Troubleshooting Commands:"

# Check NTP synchronization status on nodes
oc debug node/host01.ove.test.com -- chroot /host chrony sources -v

# Verify chrony configuration is applied
oc debug node/host01.ove.test.com -- chroot /host cat /etc/chrony.conf

# Check time synchronization status
oc debug node/host01.ove.test.com -- chroot /host timedatectl status

# Test NTP connectivity to internal server
oc debug node/host01.ove.test.com -- chroot /host chrony sourcestats

# Check for time sync issues in logs
oc debug node/host01.ove.test.com -- chroot /host journalctl -u chronyd

# Verify NTP server is accessible
ping 192.168.1.13
telnet 192.168.1.13 123
```

Storage Issues

```
echo "💾 Storage Troubleshooting Commands:"
```

```
# Check NFS connectivity from nodes
oc debug node/host01.ove.test.com -- chroot /host showmount -e 192.168.1.15

# Test NFS mount
oc debug node/host01.ove.test.com -- chroot /host mount -t nfs 192.168.1.15:/exports/nfs-ssd /mnt

# Check storage class configuration
oc get storageclass nfs-ssd -o yaml

# Check CSI driver pods
oc get pods -n nfs-csi-driver

# Check PVC status
oc get pvc --all-namespaces

# Check persistent volume status
oc get pv

# Test storage performance
oc debug node/host01.ove.test.com -- chroot /host dd if=/dev/zero of=/tmp/test bs=1M count=100
```

OpenShift Virtualization Issues

```
echo "Virtualization Troubleshooting Commands:"

# Check HyperConverged status
oc get hco -n openshift-cnv kubevirt-hyperconverged -o yaml

# Check KubeVirt status
oc get kv -n openshift-cnv kubevirt-kubevirt -o yaml

# Check virt-launcher pods
oc get pods -n openshift-cnv | grep virt-launcher

# Check node readiness for virtualization
oc get nodes -o json | jq '.items[].status.allocatable | {"kubevirt.io/kvm":
["kubevirt.io/kvm"], "tun": ["tun.network.kubevirt.io"]}'

# Check VM status
oc get vm --all-namespaces
oc get vmi --all-namespaces

# Check virtualization events
oc get events -n openshift-cnv --sort-by='.lastTimestamp'
```

```
# Verify virtualization features
oc debug node/host01.ove.test.com -- chroot /host lsmod | grep kvm
```

Operator Issues

```
echo "🔧 Operator Troubleshooting Commands:"
```

```
# Check subscription status
oc get subscription -A
```

```
# Check InstallPlan status
oc get installplan -A
```

```
# Check CSV status
oc get csv -A
```

```
# Check operator pod logs
oc logs -n openshift-cnv deployment/hco-operator
oc logs -n open-cluster-management deployment/multicloud-operators-hub-subscription
```

```
# Check catalog source status
oc get catalogsouce -n openshift-marketplace
```

```
# Check for operator events
oc get events -n openshift-marketplace --sort-by='.lastTimestamp'
```

```
# Verify operator image availability
oc get csv -A -o
jsonpath='{.items[*].spec.install.spec.deployments[*].spec.template.spec.containers[*].image}' | tr ' ' '\n' | sort -u
```

Installation Issues

```
echo "🚀 Installation Troubleshooting Commands:"
```

```
# Check bootstrap logs
ssh core@bootstrap.ove.test.com sudo journalctl -b -f -u release-image.service -u
bootkube.service
```

```
# Check installer logs
openshift-install gather bootstrap --bootstrap 192.168.1.25 --master 192.168.1.21
```

```
# Check cluster operator status
oc get clusteroperators
```

```
# Check pending CSRs
oc get csr
```

```
# Check machine configs
oc get machineconfig
oc get machineconfigpool

# Check ignition files
curl -k http://content.test.com/ignition/bootstrap.ign | jq .
```

Performance Issues

```
echo "⚡ Performance Troubleshooting Commands:"
```

```
# Check resource usage on nodes
oc adm top nodes
oc adm top pods --all-namespaces

# Check disk I/O
oc debug node/host01.ove.test.com -- chroot /host iostat -x 1

# Check network performance
oc debug node/host01.ove.test.com -- chroot /host ss -tuln

# Check memory usage
oc debug node/host01.ove.test.com -- chroot /host free -m

# Check CPU usage
oc debug node/host01.ove.test.com -- chroot /host top -b -n1
```